

70B parameter model in production

SageMaker with
tensor

parallelism across
multiple GPU

(ml.p4d.24xlarge =
8x A100s)

using LMI (Large Model
Inference) container with
DeepSpeed or vLLM

For very large models (175B+) I'd
use FSDP or Megatron-LM