



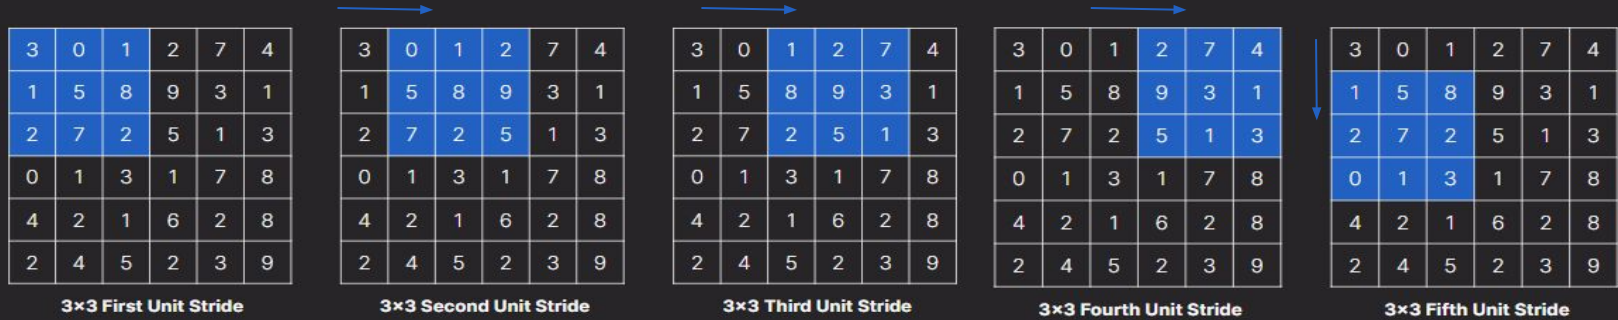
# Module 1: A Beginner's Guide to Computer Vision

## Video 5: Understanding a CNN - Striding & Padding

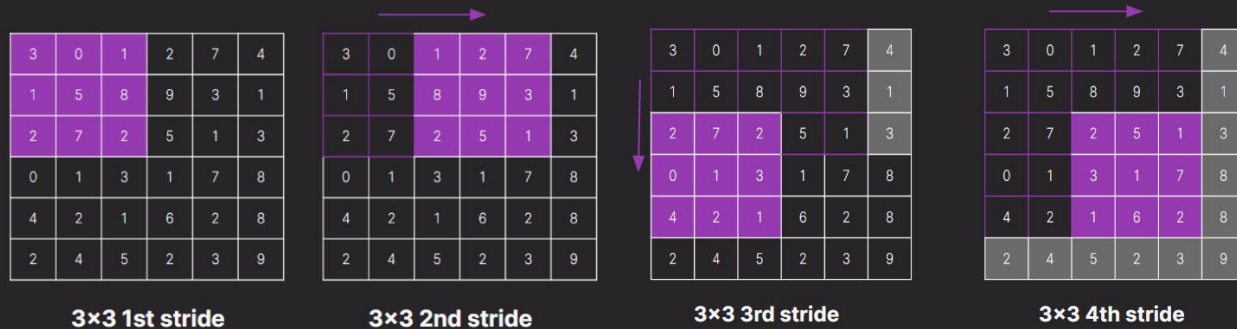
# Striding

Step size or the number of pixels shifts over the input image.

Stride = 1



Stride = 2



# Stride = 1

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

3×3 First Unit Stride

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

3×3 Second Unit Stride

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

3×3 Third Unit Stride

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

3×3 Fourth Unit Stride

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

3×3 Fifth Unit Stride

- ✓ High resolution output.
- ✓ Retains spatial information.




High computational cost.



Large output size not always ideal.

# Stride = 2



- ✓ Reduced computational load.
- ✓ Smaller output ideal for deeper networks.
- ☹ Some features are missed.
- ☹ Some loss of spatial resolution.

# The Math Behind Convolution

$$H_{out} = \left\lfloor \frac{(H_{in} - F + 2p)}{s} \right\rfloor + 1$$
$$W_{out} = \left\lfloor \frac{(W_{in} - F + 2p)}{s} \right\rfloor + 1$$

- $W_{in}$  = Input image width. Example: 6 for a 6×6 input matrix.
- $H_{in}$  = Input image height
- $F$  = Filter size (width and height are assumed to be the same)
- $p$  = Padding
- $s$  = Stride value

# The Math Behind Convolution

$$H_{out} = \left\lfloor \frac{(H_{in} - F + 2p)}{s} \right\rfloor + 1 = \frac{6 - 3 + 2 \times 0}{1} + 1 = 4$$

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

6 × 6 Matrix of the image

\*

-5	0	5
-0.5	0	0.5
-5	0	5

3×3 Filter

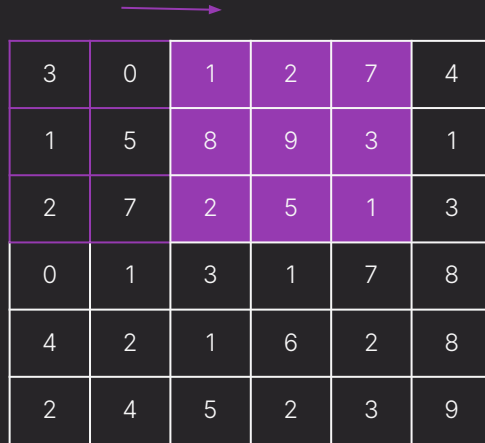
=

-6.5	2	22.5	-4
50	19	-5.5	-6
-13.5	10	2	3.5
28.5	-8	10.5	71

Convolved Matrix for the image with stride = 1  
4×4

# The Math Behind Convolution

$$\text{Output\_size} = \frac{6 - 3 + 2 \times 0}{2} + 1 = 2.5$$



3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

Stride = 2



Round down to 2

- Resulting in a 2×2 matrix

# The Math Behind Convolution

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

6×6 Matrix of the image

\*

-5	0	5
-0.5	0	0.5
-5	0	5

3×3 Filter

=

8.5	-4
28.5	71

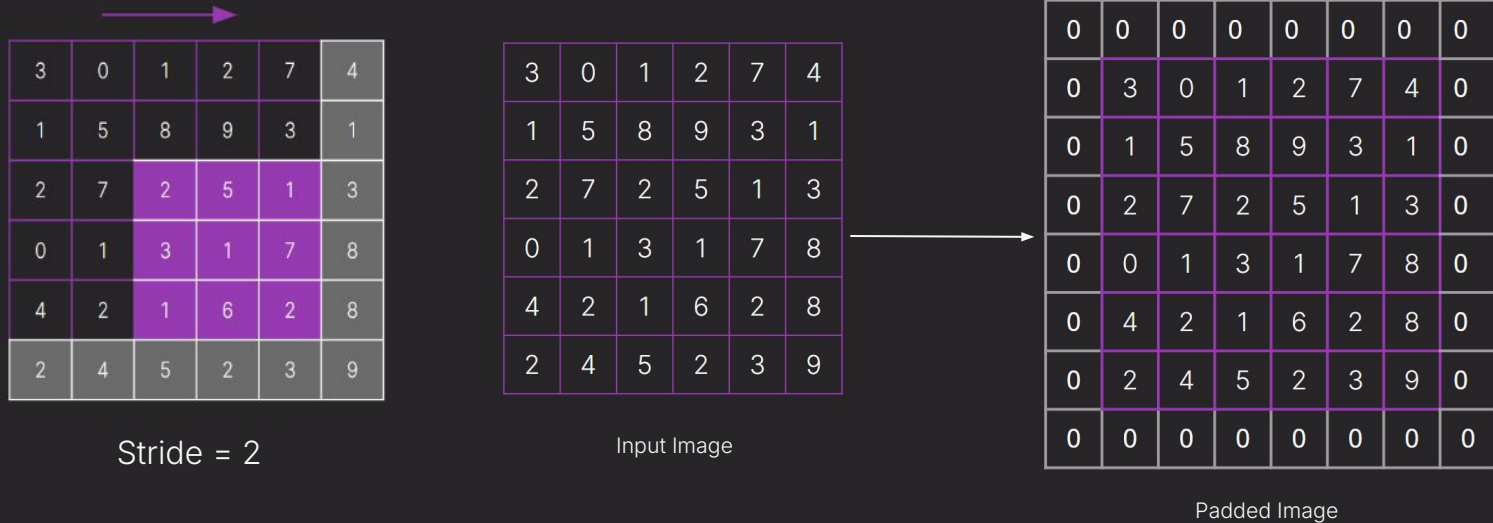
**Convolved Matrix for the Image**

(Stride = 2) 2×2



# Padding

- Padding preserves edge information in the feature map.
- Padding helps in increasing the spatial size of the output.



# Zero Padding

0	0	0	0	0	0	0	0
0	3	0	1	2	7	4	0
0	1	5	8	9	3	1	0
0	2	7	2	5	1	3	0
0	0	1	3	1	7	8	0
0	4	2	1	6	2	8	0
0	2	4	5	2	3	9	0
0	0	0	0	0	0	0	0

- **Same padding** ensures that output feature map has same dimensions as input image

# The Math Behind Convolution

$$H_{out} = \left\lfloor \frac{(H_{in} - F + 2p)}{s} \right\rfloor + 1 = \frac{6 - 3 + 2 \times 1}{1} + 1 = 6$$

0	0	0	0	0	0	0	0
0	3	0	1	2	7	4	0
0	1	5	8	9	3	1	0
0	2	7	2	5	1	3	0
0	0	1	3	1	7	8	0
0	4	2	1	6	2	8	0
0	2	4	5	2	3	9	0
0	0	0	0	0	0	0	0

\*

-5	0	5
-0.5	0	0.5
-5	0	5

3×3 Filter

=

25	34	21	-22	-39	-18.5
37.5	-6.5	2	22.5	-4	-41.5
33.5	50	19	-5.5	-6	-50.5
45.5	-13.5	10	2	3.5	-18.5
26	28.5	-8	10.5	71	-51
12	-13.5	19	4	13.5	-11.5

# Important Note



In CNNs, filter values are trainable weights.

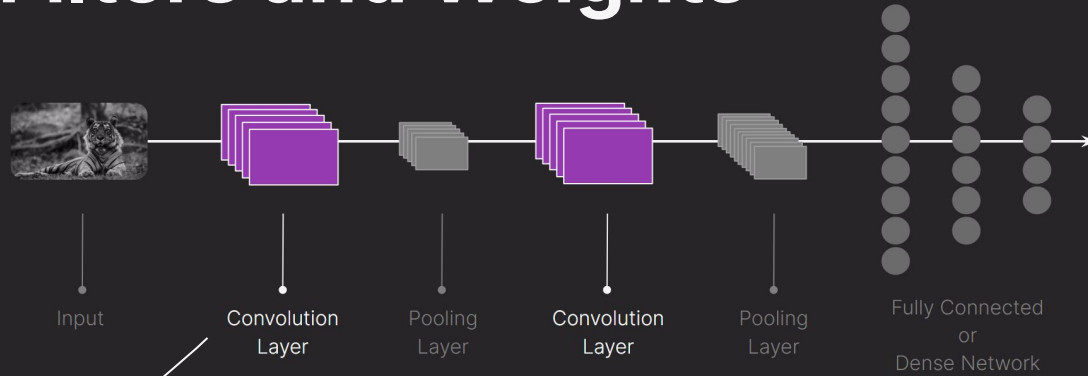


ReLU activation filter is applied to convert negative values to 0.



Size of the filters is a hyperparameter.

# Role of Filters and Weights



A curved arrow points from the first Convolution Layer to this 3x3 filter matrix:

-5	0	5
-0.5	0	0.5
-5	0	5

A 3x3 grid of weights, each labeled  $w_i$ , with a vertical column of numbers 3, 6, 9 to its right:

$w_1$	$w_2$	$w_3$	3
$w_4$	$w_5$	$w_6$	6
$w_7$	$w_8$	$w_9$	9

# Role of Activation Functions

Input

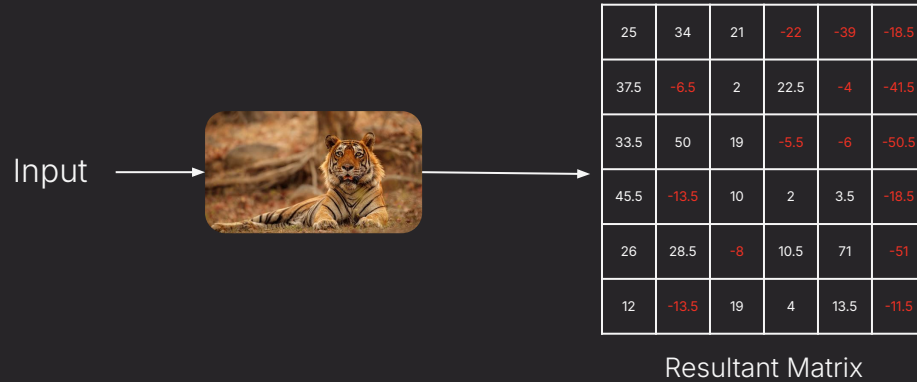


25	34	21	-22	-39	-18.5
37.5	-6.5	2	22.5	-4	-41.5
33.5	50	19	-5.5	-6	-50.5
45.5	-13.5	10	2	3.5	-18.5
26	28.5	-8	10.5	71	-51
12	-13.5	19	4	13.5	-11.5

Resultant Matrix

Apply the ReLU Activation Function

# Role of Activation Functions



- On applying the ReLU Activation Function



# Role of Activation Functions



ReLU activation filter is applied to convert negative values to 0.

Input



25	34	21	-22	-39	-18.5
37.5	-6.5	2	22.5	-4	-41.5
33.5	50	19	-5.5	-6	-50.5
45.5	-13.5	10	2	3.5	-18.5
26	28.5	-8	10.5	71	-51
12	-13.5	19	4	13.5	-11.5

Resultant Matrix

Apply ReLU

25	34	21	0	0	0
37.5	0	2	22.5	0	0
33.5	50	19	0	0	0
45.5	0	10	2	3.5	0
26	28.5	0	10.5	71	0
12	0	19	4	13.5	0

Final Matrix after ReLU  
and Bias



# Role of Filters in CNNs



Size of these filters is a hyperparameter.



Odd sizes like 3x3 are preferred.

-5	0	5
-0.5	0	0.5
-5	0	5

3×3 Matrix

1	0	0	0	-1
1	0	0	0	-1
1	0	5	0	-1
1	0	0	0	-1
1	0	0	0	-1

5×5 Matrix

0	0	0	1	0	0	0
0	0	1	2	1	0	0
0	1	3	5	3	1	0
1	2	5	8	5	2	1
0	1	3	5	3	1	0
0	0	1	2	1	0	0
0	0	0	1	0	0	0

7×7 Matrix

